

Non-Generative Learning to Rerank for MS/MS-Based Molecule Retrieval

Abstract. Retrieving molecular structures from tandem mass spectra is difficult because many candidates share mass, formula, and substructure information. Joint embedding methods rank candidates by spectrum-molecule similarity, whereas generative retrieval can introduce an autoregressive molecular decoder. We investigate an alternative design point: direct, non-generative learning to rerank a retrieved molecular list. Our framework, SpecEmbedding-Rerank, first aligns a peak-sequence spectrum encoder and a molecular graph encoder. A second-stage residual ranker then combines spectrum and candidate embeddings, explicit interaction features, the base score, and the base rank under listwise supervision. We compare a pointwise scorer without learned cross-candidate interaction with a candidate-set Transformer across three reranker initializations at a canonical first-stage checkpoint; a separate analysis repeats this design for alignment seeds 42–44. On the 17,556-query MassSpecGym test split, the canonical mean $\pm$ sample-SD Recall@1 values are 67.44 ± 0.74 and 68.58 ± 0.20 for mass-conditioned retrieval, and 73.97 ± 0.59 and 74.57 ± 0.28 for formula-conditioned retrieval, compared with fixed-base values of 47.46 and 63.10. Recall and MRR use supplied candidate files with a local exact-target-SMILES single-positive rule; the unquantified difference from reference two-dimensional InChIKey identity, which may yield multiple positives, means they are not equivalent to official evaluator outputs. Across the three alignment seeds, both learned rankers improve their corresponding base in all 12 alignment-candidate-learned-model aggregate cells for both Recall@1 and MRR, whereas the Transformer-pointwise direction varies by alignment. A fixed-alignment audit of five removal settings gives none a full-model MRR advantage in all six paired comparisons; removing the rank embedding improves Recall@1 and MRR in all six. Thus, the evidence supports supervised non-generative second-stage reranking; robust self-attention and independent component benefits remain unestablished. Neither variant generates molecules or forces the target into the test top-40 list.

Keywords: Data mining for bioinformatics · Tandem mass spectrometry · Molecule retrieval · Learning to rank · Multimodal representation learning

1 Introduction

MS/MS is widely used to identify small molecules in metabolomics, natural-product discovery, and related biomedical applications. An observed spectrum is only a partial, condition-dependent view of a molecular structure: different

molecules can yield similar fragments, while the same molecule can produce different spectra across instruments, adducts, and collision energies. Consequently, structure annotation is naturally a large-scale data mining problem in which a query spectrum must retrieve and rank a correct structure from a chemically constrained candidate pool.

MassSpecGym provides structure-disjoint data splits and standardized mass- and formula-conditioned candidate pools for this problem [3]. Recent methods mainly address the cross-modal gap between spectra and molecules. JESTR learns a spectrum-molecule joint space with contrastive multiview coding and ranks candidates directly by cosine similarity [13,21]. GLMR instead retrieves contextual candidates, generates a refined molecule with a conditional language model, and reranks candidates by molecule-molecule similarity [25].

Both approaches leave an important data mining question open. Fixed pairwise cosine scores do not directly optimize the final candidate order, while generative retrieval adds an autoregressive decoder and makes ranking depend on the quality and validity of a generated structure. We ask whether hard candidates can instead be separated directly in embedding space by a supervised second-stage ranker, and whether explicit candidate interaction adds value beyond a pointwise residual scorer.

Building on the peak-sequence spectrum architecture introduced by SpecEmbedding [23], we construct SpecEmbedding-Rerank, a two-stage retrieval and reranking framework. The first stage trains that spectrum tower from scratch and aligns it with a GINE molecular graph encoder. The second stage operates only on the first-stage top- K list. It constructs explicit spectrum-candidate interaction features, injects the original score and rank as retrieval priors, and predicts a residual correction to the base score. We instantiate this framework with pointwise and candidate-set Transformer variants.

Our contributions are:

- We formulate and evaluate direct, non-generative second-stage learning-to-rank for MS/MS-to-molecule retrieval, avoiding both fixed cosine-only ranking and autoregressive molecular decoding.
- We compare pointwise and self-attentive rankers with matched features, score heads, objectives, and optimization. On the same fixed caches, we then audit five removal settings of the full configuration.
- Using the MassSpecGym split and candidate files under our local exact-SMILES single-positive rule, both implementations improve their fixed base retriever across alignment seeds 42–44. The candidate-set Transformer advantage is not directionally uniform across these checkpoints, so the analysis does not establish a robust self-attention benefit. Moreover, within the canonical fixed-alignment audit, no removal setting gives a uniform paired MRR advantage to the full model under both candidate protocols.

AI-assistance disclosure. OpenAI Codex assisted throughout all sections with drafting and language revision, and during the project with code editing, experiment orchestration, and consistency audits. Reported observations and numerical results came from the stated software pipeline, not model generation. The

authors reviewed all AI-assisted text and code, verified the technical claims, citations, saved artifacts, and reported results, and assume full responsibility for the manuscript.

2 Related Work

2.1 Annotation Paradigms

Computational MS/MS annotation differs in retrieved objects and ranking evidence. Spectral-library systems return measured spectra and depend on library coverage and acquisition compatibility. Candidate-score correction predates our work: MetFusion combines in-silico fragmentation with spectral-library evidence [8]; MS2Query retains the top 2,000 MS2DeepScore matches and reranks them with a five-feature random forest, with optional score-threshold filtering [12]; and LC-MS2Struct combines an MS/MS scorer with chromatographic retention order to jointly rank candidates from one LC analysis [1]. Our setting instead ranks structures in query-specific MassSpecGym candidate pools from frozen spectrum-molecule representations, without spectral-library-neighbor or retention-time evidence [3].

Within structure-database retrieval, CSI:FingerID and MIST predict molecular fingerprints from spectra [7,9]. CMSSP, JESTR, and MVP instead learn joint spectrum-molecule spaces by contrastive alignment [5,13,6]; MVP aligns molecular graphs and fingerprints with individual and consensus spectra. These methods ultimately order candidates using a predicted-fingerprint score or cross-modal similarity. SpecEmbedding addresses spectrum-spectrum representation learning [23]. We reuse only its peak-sequence architecture, train it from scratch against molecular graphs, and study the complementary second-stage problem of learning the order within an already retrieved hard-candidate list.

Simulation and generation use different intermediate representations. MassFormer predicts tandem mass spectra from molecular graphs [24]; MSNovelist generates structures from MS/MS-derived fingerprints [20]. GLMR aligns spectra and molecules, conditions a molecular decoder on retrieved candidates, and orders those candidates by similarity to the generated molecule [25]. Our method does not decode a structure: it remains a closed-set reranker and cannot return a molecule outside the first-stage list, but directly optimizes the scores of available candidates.

2.2 Learning to Rank Candidate Sets

Learning-to-rank objectives define supervision over preferences or complete lists rather than treating similarity sorting as a fixed post-processing step. RankNet learns pairwise preferences, whereas ListNet defines a probability model over a ranked list [2,4]. Ranking functions can also exploit relations among competing items. Set Transformer provides attention-based blocks for set-structured

Table 1. Mechanistic comparison with the closest MassSpecGym approaches.

Model	Ranking signal	Candidate interaction	Generates	Training stages
JESTR	Spectrum-candidate cosine	None at inference	No	InfoNCE plus formula-candidate regularizer
GLMR	Generated-candidate cosine	Generator cross-attention	Yes	Dual-path InfoNCE plus generation
Ours	Residual score from pair features and retrieval priors	Optional candidate self-attention	No	Cross-modal InfoNCE plus listwise/pairwise reranking

inputs [15], and SetRank jointly models cross-document interactions while preserving permutation equivariance [17]. We transfer these principles to spectrum-molecule retrieval using cached cross-modal embeddings, explicit pair interactions, and first-stage score and rank priors. Our contribution is not self-attention or listwise learning by themselves, but a controlled non-generative second-stage formulation and evaluation for MS/MS retrieval.

2.3 Positioning of Our Method

Table 1 isolates the ranking mechanisms of the closest MassSpecGym approaches. In light of the precedents above, our contribution is not reranking itself, but a query-specific, non-generative second-stage formulation over frozen cross-modal representations. Residual scoring and candidate self-attention are evaluated design choices rather than assumed sources of gain.

3 Problem Formulation

Let s be an MS/MS spectrum, m^+ its correct molecule, and $\mathcal{D}_s$ a query-specific mass- or formula-conditioned candidate pool. A base retriever assigns

$$b(s, m) = \text{sim}(f_s(s), f_m(m)) \quad (1)$$

to every $m \in \mathcal{D}_s$ and returns the top- K list $\mathcal{C}_s = (m_1, \dots, m_K)$. The reranker learns scores $g(s, m_i, \mathcal{C}_s)$ such that m^+ is ordered as high as possible. Because it cannot return a molecule outside $\mathcal{C}_s$, end-to-end Recall@ k is upper-bounded by

$$U_K = \frac{1}{N} \sum_{n=1}^N \mathbb{I}[m_n^+ \in \mathcal{C}_{s_n}]. \quad (2)$$

Define the within-list top- k accuracy conditional on first-stage coverage as




Fig. 1. Overview of SpecEmbedding-Rerank. Panel 1 aligns the spectrum Transformer and molecular GINE. Panel 2 uses their frozen, cached embeddings to produce the top-40 list. Panel 3 applies shared pair features with either pointwise scoring or candidate-set self-attention and adds a residual to the base score. Dashed arrows denote training-only operations. Positives are never forced into the truncated top-40 list at validation/test, and no molecule is generated.

$$A_{k|K} = \frac{\sum_{n=1}^N \mathbb{I}[m_n^+ \in \mathcal{C}_{s_n}, \text{rank}_g(m_n^+) \leq k]}{\sum_{n=1}^N \mathbb{I}[m_n^+ \in \mathcal{C}_{s_n}]}, \quad \text{Recall@}k = U_K A_{k|K} \leq U_K. \quad (3)$$

A reranker applied to a fixed cache can therefore change only $A_{k|K}$, not U_K . We report metrics over all queries, assigning zero credit when the positive is absent rather than conditioning the primary metric on successful first-stage retrieval.

4 Method

4.1 Overview

The framework has three steps: cross-modal alignment, top- K retrieval, and supervised residual reranking. Spectrum and molecule encoders are frozen while the reranker is trained, so their embeddings can be computed once and reused.

4.2 Cross-Modal Base Retriever

The spectrum encoder follows SpecEmbedding’s peak-sequence architecture [23]: it processes peak m/z values and intensities with sinusoidal m/z features and a Transformer. Each spectrum is represented by at most 100 tokens: a precursor- m/z token with intensity 2, followed by up to 99 fragment peaks selected by

intensity. Selected fragments retain their source-array order, and their intensities are normalized by the largest selected fragment intensity. Only these m/z -intensity tokens enter the spectrum encoder; adduct, instrument, and collision-energy metadata are not explicit inputs. We initialize this tower from scratch and train it against molecules rather than spectra. The molecular encoder is a GINE network over atom and bond features [11]. Mean-pooled graph features are augmented with explicit graph-size features. Two projection heads map the encoders to a common d -dimensional space:

$$\mathbf{z}_s = \frac{p_s(f_s(s))}{\|p_s(f_s(s))\|_2}, \quad \mathbf{z}_m = \frac{p_m(f_m(m))}{\|p_m(f_m(m))\|_2}. \quad (4)$$

We use a bidirectional cross-modal contrastive loss that permits multiple positives. Spectra and molecules sharing a molecular label form positives, which avoids treating replicate spectra as false negatives. For a minibatch of size B , let $P(i)$ be the indices sharing the label of item i . With a learned inverse-temperature scale a , the base logits are

$$\ell_{ij} = a \mathbf{z}_{s_i}^\top \mathbf{z}_{m_j}. \quad (5)$$

The spectrum-to-molecule term is

$$\mathcal{L}_{s \rightarrow m} = -\frac{1}{B} \sum_{i=1}^B \frac{1}{|P(i)|} \sum_{j \in P(i)} \log \frac{\exp(\ell_{ij})}{\sum_{k=1}^B \exp(\ell_{ik})}, \quad (6)$$

and the molecule-to-spectrum term applies the same expression to $\ell^\top$. The alignment loss is their mean. The implementation parameterizes $a = \exp(u)$ and caps it at 100. At inference, the normalized dot product $b_i = \mathbf{z}_s^\top \mathbf{z}_{m_i}$ is the base score.

4.3 Candidate Construction and Training Protocol

MassSpecGym supplies one candidate pool per query. The base retriever scores the pool and keeps its top $K = 40$ molecules. If a training target is absent from the retrieved top-40, it replaces the last molecule while its full-pool base rank is cached. That rank is one plus the number of strictly higher-scoring candidates; ties with the target do not count as higher. The rank embedding receives $\bar{r}_i = \min\{\max(r_i, 0), 512\}$, with index 0 reserved for padding. Training also shuffles candidate order while moving each candidate’s embedding, score, and rank together to guard against cache-order artifacts.

Scored pools are exact-SMILES-deduplicated unions of the labeled target and supplied candidates. A static audit found the target already present in all 32,010 frozen pools, so no pool was enlarged. Post-retrieval positive forcing is training-only; validation and test perform no top-40 replacement. A target absent from the retrieved top-40 remains absent and lowers the upper bound and all end-to-end metrics. No test record, target, or label is used for optimization or checkpoint selection.

4.4 Residual Reranker Variants

For candidate m_i with base rank r_i , we construct

$$\mathbf{x}_i = [\mathbf{z}_s, \mathbf{z}_{m_i}, \mathbf{z}_s \odot \mathbf{z}_{m_i}, |\mathbf{z}_s - \mathbf{z}_{m_i}|, b_i, \mathbf{e}(\bar{r}_i)], \quad (7)$$

where $\mathbf{e}(\bar{r}_i)$ is a learned rank embedding. The direct embeddings retain modality-specific evidence, while product and absolute-difference terms make symmetric matching cues explicit. An MLP projects each vector:

$$\mathbf{h}_i = \text{MLP}(\mathbf{x}_i). \quad (8)$$

Both variants receive only the frozen spectrum embedding, frozen candidate embeddings, cosine base scores, clipped base-rank embeddings, and the candidate padding mask. Raw peaks, SMILES tokens, precursor mass or formula, and acquisition metadata are not passed directly to the reranker. Mass and formula constraints affect candidate-pool construction only; no explicit candidate-candidate structural-similarity feature is provided.

The candidate-set variant uses a Transformer encoder without candidate positional encodings to exchange information across candidates [22]:

$$[\tilde{\mathbf{h}}_1, \dots, \tilde{\mathbf{h}}_K] = \text{Transformer}([\mathbf{h}_1, \dots, \mathbf{h}_K]). \quad (9)$$

Padding masks support variable-sized lists. At deterministic inference, jointly permuting the candidates and their attached embeddings, scores, and ranks permutes the output scores in the same way; the base rank is a candidate-attached retrieval prior, not a tensor-slot position encoding [15,17]. The pointwise control sets $\tilde{\mathbf{h}}_i = \mathbf{h}_i$, retaining the same features, score head, residual path, and training losses while removing learned cross-candidate interaction after the retrieval priors have been constructed. It is not parameter- or capacity-matched because it removes the two Transformer layers: the pointwise and candidate-set variants contain exactly 1,478,690 and 7,783,458 trainable reranker parameters, respectively. A score head predicts a correction Δ_i , and the final score is residual:

$$q_i = \alpha b_i + \Delta_i, \quad (10)$$

where α is initialized to 1.0 and learned without a sign constraint. The residual path provides a direct shortcut from the base score; it does not guarantee preservation of the base order. Moreover, because b_i also enters $\mathbf{x}_i$ and Δ_i can learn additional score-dependent effects, α alone is not a mechanistically interpretable measure of reliance on the base retriever.

4.5 Learning Objective

For the labeled queries $\mathcal{B}_L$ in a minibatch, let V_n be the valid candidate set and y_n the positive index. The primary objective averages listwise cross-entropy over labeled queries:

$$\mathcal{L}_{\text{list}} = -\frac{1}{|\mathcal{B}_L|} \sum_{n \in \mathcal{B}_L} \log \frac{\exp(q_{ny_n})}{\sum_{i \in V_n} \exp(q_{ni})}. \quad (11)$$

We also contrast each positive with every other valid candidate retrieved for that query. The implementation averages over all valid positive-negative pairs in the minibatch:

$$\mathcal{L}_{\text{pair}} = \frac{\sum_{n \in \mathcal{B}_L} \sum_{i \in V_n \setminus \{y_n\}} \log(1 + \exp(q_{ni} - q_{ny_n} + \gamma))}{\sum_{n \in \mathcal{B}_L} (|V_n| - 1)}. \quad (12)$$

The complete objective is

$$\mathcal{L} = \mathcal{L}_{\text{list}} + \lambda \mathcal{L}_{\text{pair}}. \quad (13)$$

The listwise term is a top-one likelihood: it separates the positive from the valid list but does not prescribe a total order among negatives. Both terms are invariant to adding the same query-specific constant to every q_i . Thus, the raw outputs are relative scores, not probabilities; probability calibration is neither performed nor evaluated.

For L Transformer layers and hidden width h , the set-aware variant adds $O(L(Kh^2 + K^2h))$ computation and $O(LK^2)$ attention memory over the pointwise scorer. Neither variant decodes molecules token by token. An empirical latency comparison with GLMR is not reported, so we make no measured efficiency claim.

5 Experiments

5.1 Dataset and Candidate Pools

We use the official MassSpecGym split [3]: 194,119 training spectra, 19,429 validation spectra, and 17,556 test spectra. These counts total 231,104 spectra; the benchmark reports approximately 29,000 unique molecular structures. Its structure-disjoint split groups similar molecules and makes test structures substantially different from training structures.

An exact audit of spectrum-record signatures (peak arrays, precursor m/z , adduct, instrument, and collision energy) found no target-molecule duplicates across splits, but six validation and three test signatures exactly matched training signatures. In all nine matched pairs, the target molecules differ and each evaluation query is marked `simulation_challenge=true`. We report these cases as cross-split spectrum-record overlap. In the canonical pipeline reported below, all validation-dependent model selection excludes the six affected validation spectra, leaving 19,423 validation queries for the rerankers; alignment validation keys decrease from 3,386 to 3,384 molecular entries. We retain and evaluate all 17,556 official test queries.

MassSpecGym supplies mass- and formula-conditioned candidate settings. The mass pool uses a precursor-mass tolerance; the formula pool assumes a

known formula and is not unrestricted annotation. Frozen pools contain at most 256 candidates. Before truncation, their test-query means are 253.88 and 165.75, respectively; 1 mass and 3,159 formula pools contain fewer than 40 candidates. After top-40 truncation, the means are 40.00 and 35.73.

Our cache and evaluator assign one positive by exact target-SMILES equality. In official code commit `4f501b3`,¹ the reference loader instead defaults to two-dimensional InChIKey equivalence and may label multiple positives. Our Recall and MRR therefore use a local exact-SMILES protocol on the supplied candidate files; the effect of the identity-rule difference is unquantified.

5.2 Implementation Details

The aligned model outputs 512-dimensional embeddings. Its spectrum Transformer has four layers, 16 heads, and width 512; the molecular GINE has four layers, hidden width 128, and dropout 0.2. No published SpecEmbedding checkpoint is loaded, so the stage that would freeze a pretrained spectrum tower is skipped; both towers are optimized end to end on the MassSpecGym training split. Alignment uses AdamW with batch size 128, learning rate and weight decay 10^{-4} , cosine scheduling, gradient-norm clipping at 1, at most 100 epochs, and patience 5. Each epoch makes ten draws per molecular label, randomly selecting an associated spectrum. Spectrum and graph augmentation are independently enabled with probability 0.5. The former removes a floor-rounded 20% of fragment peaks below normalized intensity 0.3 and jitters intensities by $\pm 15\%$; the latter applies 10% node and 10% directed edge-entry dropout. Validation disables augmentation but retains replicate-spectrum sampling under the fixed random stream. The canonical seed-42 run applies the validation exclusion above, selects epoch 16, and stops at epoch 21.

The reranker has hidden width 512, a 32-dimensional rank embedding, and dropout 0.1. The set-aware variant has two Transformer layers with eight attention heads; the pointwise control removes both layers. We use AdamW with learning rate 10^{-4} , batch size 16, at most 30 epochs, and early stopping on validation MRR with patience 5. The pairwise weight and margin are $\lambda = 0.2$ and $\gamma = 0.1$. For the canonical controlled analysis, the alignment-seed-42 checkpoint and its overlap-clean caches are fixed while each reranker variant is trained with seeds 42, 43, and 44. These sample standard deviations measure reranker variation conditional on that first stage. A separate cross-alignment sensitivity analysis applies the same reranker seeds to alignment seeds 42, 43, and 44 and first aggregates reranker seeds within each alignment. Its cross-alignment unit is therefore the alignment checkpoint ($n = 3$), not nine flattened runs. The component audit remains limited to the canonical alignment-seed-42 caches and reranker seeds. Its five removals respectively disable only the residual shortcut while retaining the base score as MLP input; jointly disable the base-score feature and shortcut; zero the rank embedding while preserving input width; remove product and absolute-difference features while retaining self-attention;

¹ MassSpecGym reference loader and label transform, accessed August 1, 2026.

or set $\lambda = 0$ for listwise-only training. All models are selected without test-set access and evaluated on one NVIDIA RTX 4090.

Before training-only positive forcing, top-40 base recall is 93.78% for mass and 90.42% for formula on the training split. Training lists are then completed to 100% labeled coverage as described above. No target is forced into the test top-40 lists, whose corresponding upper bounds are 83.88% and 89.26%.

5.3 Baselines and Metrics

The controlled baselines are the same aligned retriever without the second-stage ranker and a pointwise reranker that removes candidate self-attention while retaining all other inputs and objectives, but not the same parameter count. We separately show JESTR and GLMR values from Table 1 of the GLMR paper [25]. These external numbers are not a controlled comparison: JESTR and GLMR were not independently reproduced in our codebase, and their retrievers, candidate handling, identity/canonicalization rules, and MCES implementations are not matched to our local exact-SMILES single-positive protocol.

We report Recall@1, Recall@5, Recall@20, and mean reciprocal rank (MRR) over all test queries. For every query, including those whose target is absent from the retrieved top-40, let $\hat{m}_n^{(1)}$ be the top-ranked candidate. We report

$$\text{MCES@1} = \frac{1}{N} \sum_{n=1}^N \tilde{d}_{15}(\hat{m}_n^{(1)}, m_n^+), \quad N = 17,556. \quad (14)$$

Here $\tilde{d}_{15}$ is the output of `myopic-MCES` with threshold 15 and the stronger-bound option enabled [14]. It is exact for sufficiently close pairs; for distant pairs it may return the threshold or a computed lower bound rather than the unconstrained exact MCES distance. Thus, the reported quantity is thresholded myopic MCES, and lower is better. Exact SMILES-string matches are assigned zero. The canonical runs selected `PULP_CBC_CMD`, and no MCES worker failures were logged. Recall and MRR uncertainty is computed over three reranker seeds with a fixed base checkpoint in the canonical audit. The cross-alignment summary first averages reranker seeds within each checkpoint and then uses the three alignment-level paired estimates; neither sample standard deviation is a confidence interval. MCES@1 is computed for the overlap-clean base and representative seed-42 Transformer checkpoints; these single-checkpoint values do not include seed-level uncertainty.

5.4 Main Results

In the representative seed-42 run, the candidate-set Transformer adds 21.28 percentage points in mass-conditioned Recall@1 and 11.65 points in formula-conditioned Recall@1. The corresponding MRR gains are 17.84 and 9.36 points. Improvements at Recall@20 are smaller because the first-stage top-40 upper bound is already the dominant constraint. Relative to the corresponding fixed

Table 2. Representative local MassSpecGym results. The base uses alignment seed 42, and the learned row uses reranker seed 42. Recall and MRR are percentages under our exact-target-SMILES single-positive rule on supplied candidate files; they are not official-evaluator-equivalent. MCES@1 is the local thresholded myopic-MCES output (lower is better).

Pool	Method	R@1	R@5	R@20	MRR	MCES@1
Mass	SpecMolAlign (align. seed 42)	47.46	63.51	77.76	55.03	15.37
	Candidate-set Transformer (rerank seed 42)	68.74	77.77	82.18	72.87	7.71
Formula	SpecMolAlign (align. seed 42)	63.10	75.79	84.93	68.86	5.44
	Candidate-set Transformer (rerank seed 42)	74.75	82.25	86.89	78.22	3.09

Table 3. External cross-paper context reported by Zhang et al. [25]. These values were not independently reproduced and are not directly comparable to Table 2; retrievers, candidate and identity processing, and MCES implementations may differ.

Pool	Method	R@1	R@5	R@20	MRR	MCES@1
Mass	JESTR (reported; not reproduced)	17.62	40.36	64.76	29.12	15.82
	GLMR (reported; not reproduced)	64.17	72.96	78.78	67.82	11.14
Formula	JESTR (reported; not reproduced)	11.77	33.26	61.01	22.83	11.73
	GLMR (reported; not reproduced)	68.48	78.09	84.22	72.47	5.05

overlap-clean base, the same seed-42 Transformer reduces MCES@1 from 15.3681 to 7.7057 for mass candidates (49.86%) and from 5.4389 to 3.0913 for formula candidates (43.16%).

Table 3 records the values reported in the GLMR paper without ranking those external pipelines against ours. A matched comparison would require rerunning the baselines with the same retriever, candidate lists, identity rule, and metric implementation. The displayed external ordering therefore supports no relative-performance ranking or method attribution.

5.5 Cross-Alignment and Controlled Component Evidence

Table 4 preserves the alignment hierarchy: reranker seeds are first averaged within each first-stage checkpoint. Both learned rankers improve the alignment-matched base in Recall@1 and MRR in all 12 alignment-candidate-learned-model aggregate cells. In contrast, the Transformer-pointwise difference is positive for only two of three alignments for each candidate setting and metric. Across the three alignment-level paired estimates, the mean $\pm$ sample-SD differences are 0.35 ± 0.75 Recall@1 points and 0.22 ± 0.56 MRR points for mass, and 0.11 ± 0.47 and 0.13 ± 0.29 points for formula. These are descriptive summaries, not confidence intervals or significance tests. The three checkpoints share the analyzed parameter hash but were produced from two source commits; per-alignment provenance is retained in the analysis artifacts. Thus, these internal

Table 4. Alignment-stratified internal results under our local exact-target-SMILES single-positive, MassSpecGym supplied-candidate top-40 protocol, which is not official-evaluator-equivalent. Learned entries are means over reranker seeds 42–44 within each alignment. The three checkpoints come from two source commits and provide descriptive sensitivity evidence, not confidence intervals, significance tests, or general end-to-end random-seed stability. Recall@1 and MRR are percentages.

Alignment	Pool	Base R@1	Point. R@1	Trans. R@1	Base MRR	Point. MRR	Trans. MRR
42	Mass	47.46	67.44	68.58	55.03	71.79	72.60
42	Formula	63.10	73.97	74.57	68.86	77.64	78.02
43	Mass	35.92	66.92	67.19	44.33	70.42	70.58
43	Formula	45.08	69.64	69.30	51.85	72.14	71.96
44	Mass	23.91	57.78	57.42	32.26	61.22	60.91
44	Formula	28.73	59.76	59.82	36.87	63.14	63.34

Table 5. Controlled comparison on fixed alignment-seed-42 caches. The base row is fixed; learned rows report mean $\pm$ sample standard deviation over reranker seeds 42–44. Recall@1 and MRR are percentages. Removal rows are separately trained audits of the predefined full configuration, not post-hoc model-selection candidates. All values use the local exact-SMILES single-positive protocol and are not official-evaluator-equivalent.

Method	Mass		Formula	
	R@1	MRR	R@1	MRR
Base	47.46	55.03	63.10	68.86
Pointwise	67.44 $\pm$ 0.74	71.79 $\pm$ 0.66	73.97 $\pm$ 0.59	77.64 $\pm$ 0.49
Candidate-set Transformer (full)	68.58 $\pm$ 0.20	72.60 $\pm$ 0.25	74.57 $\pm$ 0.28	78.02 $\pm$ 0.26
No residual shortcut	68.16 $\pm$ 0.06	72.33 $\pm$ 0.06	75.01 $\pm$ 0.54	78.38 $\pm$ 0.42
No base-score	68.07 $\pm$ 0.20	72.28 $\pm$ 0.18	74.39 $\pm$ 0.55	77.93 $\pm$ 0.38
input/shortcut				
No rank embedding	70.12 $\pm$ 0.30	73.75 $\pm$ 0.28	76.62 $\pm$ 0.69	79.38 $\pm$ 0.58
Product/absolute-difference removed	68.10 $\pm$ 0.85	72.29 $\pm$ 0.66	75.76 $\pm$ 0.55	78.90 $\pm$ 0.38
Listwise only ($\lambda = 0$)	68.91 $\pm$ 0.76	72.86 $\pm$ 0.46	74.24 $\pm$ 0.25	77.71 $\pm$ 0.25

results support supervised residual reranking but do not establish a checkpoint-robust independent benefit of candidate self-attention.

On the canonical alignment-seed-42 checkpoint in Table 5, the candidate-set Transformer exceeds the pointwise model by 1.14/0.60 Recall@1 points and 0.81/0.38 MRR points for mass/formula candidates, respectively. All six same-seed Recall@1 and MRR differences favor the Transformer within this checkpoint. Because it also has 6.30M more trainable parameters, even this within-checkpoint comparison does not isolate candidate interaction from increased capacity; the cross-alignment direction changes further preclude a checkpoint-robust self-attention claim.

The separately trained removal rows audit the full configuration on the same fixed alignment. No removal setting gives the full model higher MRR in all six paired comparisons. Removing the rank embedding improves both reported metrics in all six; its mean Recall@1/MRR gains are 1.54/1.15 mass points and 2.05/1.36 formula points. The residual shortcut, product/absolute-difference fea-

tures, and pairwise loss have candidate-dependent directions. The full model’s mean MRR lead over the joint base-score-pathway removal is only 0.32/0.08 points, with mixed seed signs. These are descriptive comparisons, not significance tests; the predefined full configuration remains the reporting reference.

On the earlier checkpoint, a sensitivity check excluded the three affected test queries, leaving 17,553. Across both candidate settings and rerankers, mean Recall@1 decreased by 0.0045–0.0055 points (maximum seed-level change: 0.0056), so the aggregate effect is negligible; canonical tables retain all test queries.

The clean and earlier pipelines select the same alignment epochs (best 16, stop 21), but reranker best/stop epochs differ in 9 of 12 configurations. Because alignment was retrained and its test metrics changed, this measures pipeline sensitivity rather than the causal effect of six validation queries.

We do not report candidate-order or full top- K ablations. We also do not report peak memory, cache construction time, reranker latency, or end-to-end latency.

6 Discussion

Equations 2–3 separate first-stage coverage from within-list ordering. For $k = 1$ and $K = 40$, the defined conditional quantity $A_{1|40}$ equals $R@1/U_{40}$. Using the displayed three-seed means, it increases from 56.58% for the base retriever to 80.40% for the pointwise model and 81.76% for the candidate-set Transformer under mass conditioning. The corresponding formula-conditioned rates are 70.69%, 82.87%, and 83.54%.

Another view is the fraction of attainable top-1 headroom recovered,

$$H = \frac{R@1_{\text{rerank}} - R@1_{\text{base}}}{U_{40} - R@1_{\text{base}}}. \quad (15)$$

The pointwise and candidate-set models recover 54.86% and 57.99% of this headroom for mass candidates, and 41.55% and 43.85% for formula candidates. The pointwise gain is also 94.60% and 94.77% of the total candidate-set gain over the base model in the two settings. Thus, most of the observed improvement is already achieved by the supervised pointwise second-stage ranker as configured; the candidate-set variant adds only a modest increment. At Recall@20, the candidate-set three-seed means of 82.20% and 86.77% are only 1.68 and 2.49 points below U_{40} ; first-stage coverage is therefore the dominant remaining limitation for that metric under this top-40 protocol. The removals likewise support only a framework-level interpretation: none gives the full model uniform paired MRR support, whereas rank removal improves both metrics in every pair. The predefined full configuration remains the reporting reference; these algebraic and removal comparisons are not causal estimates of individual components. Across alignment seeds 42–44, both learned variants retain a large descriptive margin over their corresponding base, despite substantial first-stage metric variation. This strengthens the framework-level observation within the three audited runs, while the mixed Transformer–pointwise directions reinforce that most of the gain should not be attributed to candidate self-attention.

7 Limitations and Responsible Evaluation

Three-seed sample standard deviations and paired-win counts are descriptive, not confidence intervals or significance tests. The canonical table and all component removals fix alignment seed 42. The separate cross-alignment analysis covers only three internal checkpoints, produced by two source commits, and first aggregates reranker seeds within each alignment; it is descriptive sensitivity evidence rather than a confidence interval or general stability claim. The component audit therefore establishes neither cross-alignment component stability nor causal effects: no-base-score jointly removes its MLP input and shortcut, product/absolute-difference removal retains self-attention, and the rank-free variant was neither tested across alignments nor selected post hoc. The candidate-set model also has 5.26 times the pointwise model’s trainable parameters (6,304,768 more; +426%), confounding interaction with capacity. External JESTR and GLMR rows were not reproduced under identical software and candidate handling; we claim neither definitive state of the art nor an independent component benefit.

Training-only positive replacement changes the list distribution without addressing $1 - U_K$. Because frozen full pools contain the exact target, pool-level evaluation is closed-set even when top-40 retrieval misses it. The impact of exact-SMILES rather than reference two-dimensional InChIKey identity is unquantified, so Recall/MRR are explicitly local-protocol results rather than official-evaluator-equivalent values. The exact-record audit cannot exclude near duplicates or correlated provenance. Model selection excludes the six audited training-validation overlaps; all test queries remain. The post-hoc validation environment is now version-locked and was rebuilt from explicit Conda/PyPI locks, but historical training runs did not preserve a per-package environment record; therefore we do not claim bitwise retraining equivalence. A 2026 audit preprint reports further MassSpecGym leakage, shortcut, implementation, and metric risks and releases v1.5 [16]; our scores are artifact- and protocol-specific, not proven v1.5-equivalent.

Alignment sampling approximately balances molecular labels, whereas metrics count spectra, giving replicate-rich molecules more weight. Candidate-pool Recall measures prioritization, not confirmed identification; confidence frameworks and work on spectral-library absences distinguish these notions [19,10]. Conformal retrieval can produce spectrum-specific sets with marginal coverage under exchangeability and target-in-pool assumptions [18], but such guarantees do not automatically cover a target missing from the retrieved top-40. We perform neither probability calibration nor abstention; orthogonal experimental evidence remains necessary.

Finally, formula conditioning presumes a known formula. The spectrum tower uses selected peaks and precursor m/z , but no explicit adduct, instrument, or collision-energy inputs. Robustness to acquisition shift, other retrievers and embeddings, identity rules, pool construction, larger lists, other benchmarks, and external libraries is untested. The closed-set method cannot propose structures

outside the supplied pool despite avoiding autoregressive generation-validity failures.

8 Conclusion

We presented SpecEmbedding-Rerank for supervised, non-generative second-stage reranking of MS/MS candidates. Recall and MRR use MassSpecGym-supplied candidate files with a local exact-target-SMILES single-positive rule. Because the reference loader defaults to two-dimensional InChIKey equivalence and may yield multiple positives, the unquantified difference means these results are not official-evaluator-equivalent. Within this protocol, both learned variants improve their alignment-matched base in both metrics in all 12 aggregate cells across alignment seeds 42–44. The candidate-set Transformer has a modest canonical-checkpoint advantage over the pointwise scorer, but its direction is positive in only two of three alignments for each candidate setting and metric. None of five removal settings gives the full model uniform paired MRR support within the canonical alignment, while rank removal improves both reported metrics in all six fixed-alignment pairs. The strongest supported conclusion therefore concerns the reranking framework, not an independent component benefit. The separately displayed JESTR/GLMR values were not reproduced and are not directly comparable; matched external baselines and measured efficiency remain outside the evidence.

References

1. Bach, E., Schymanski, E.L., Rousu, J.: Joint structural annotation of small molecules using liquid chromatography retention order and tandem mass spectrometry data. *Nature Machine Intelligence* **4**, 1224–1237 (2022). <https://doi.org/10.1038/s42256-022-00577-2>
2. Burges, C., Shaked, T., Renshaw, E., Lazier, A., Deeds, M., Hamilton, N., Hullender, G.: Learning to rank using gradient descent. In: *Proceedings of the 22nd International Conference on Machine Learning*. pp. 89–96 (2005). <https://doi.org/10.1145/1102351.1102363>
3. Bushuiev, R., Bushuiev, A., de Jonge, N.F., Young, A., Kretschmer, F., Samusevich, R., Heirman, J., Wang, F., Zhang, L., Dührkop, K., et al.: MassSpecGym: A benchmark for the discovery and identification of molecules. In: *Advances in Neural Information Processing Systems*. vol. 37, pp. 110010–110027 (2024). <https://doi.org/10.52202/079017-3491>
4. Cao, Z., Qin, T., Liu, T.Y., Tsai, M.F., Li, H.: Learning to rank: From pairwise approach to listwise approach. In: *Proceedings of the 24th International Conference on Machine Learning*. pp. 129–136 (2007). <https://doi.org/10.1145/1273496.1273513>
5. Chen, L., Xia, B., Wang, Y., Huang, X., Gu, Y., Wu, W., Zhou, Y.: CMSSP: A contrastive mass spectra-structure pretraining model for metabolite identification. *Analytical Chemistry* **96**(42), 16871–16881 (2024). <https://doi.org/10.1021/acs.analchem.4c03724>

6. Chen, Y.Z., Hassoun, S.: Learning from all views: A multiview contrastive framework for metabolite annotation. *Analytical Chemistry* **98**(9), 6598–6606 (2026). <https://doi.org/10.1021/acs.analchem.5c05675>
7. Dührkop, K., Shen, H., Meusel, M., Rousu, J., Böcker, S.: Searching molecular structure databases with tandem mass spectra using CSI:FingerID. *Proceedings of the National Academy of Sciences* **112**(41), 12580–12585 (2015). <https://doi.org/10.1073/pnas.1509788112>
8. Gerlich, M., Neumann, S.: MetFusion: Integration of compound identification strategies. *Journal of Mass Spectrometry* **48**(3), 291–298 (2013). <https://doi.org/10.1002/jms.3123>
9. Goldman, S., Wohlwend, J., Stražar, M., Haroush, G., Xavier, R.J., Coley, C.W.: Annotating metabolite mass spectra with domain-inspired chemical formula transformers. *Nature Machine Intelligence* **5**(9), 965–979 (2023). <https://doi.org/10.1038/s42256-023-00708-3>
10. Hoffmann, M.A., Nothias, L.F., Ludwig, M., Fleischauer, M., Gentry, E.C., Witting, M., Dorrestein, P.C., Dührkop, K., Böcker, S.: High-confidence structural annotation of metabolites absent from spectral libraries. *Nature Biotechnology* **40**(3), 411–421 (2022). <https://doi.org/10.1038/s41587-021-01045-9>
11. Hu, W., Liu, B., Gomes, J., Zitnik, M., Liang, P., Pande, V., Leskovec, J.: Strategies for pre-training graph neural networks. In: *International Conference on Learning Representations* (2020)
12. de Jonge, N.F., Louwen, J.J.R., Chekmeneva, E., Camuzeaux, S., Vermeir, F.J., Jansen, R.S., Huber, F., van der Hooft, J.J.J.: MS2Query: Reliable and scalable MS2 mass spectra-based analogue search. *Nature Communications* **14**(1), 1752 (2023). <https://doi.org/10.1038/s41467-023-37446-4>
13. Kalia, A., Chen, Y.Z., Krishnan, D., Hassoun, S.: JESTR: Joint embedding space technique for ranking candidate molecules for the annotation of untargeted metabolomics data. *Bioinformatics* **41**(7), btaf354 (2025). <https://doi.org/10.1093/bioinformatics/btaf354>
14. Kretschmer, F., Seipp, J., Ludwig, M., Klau, G.W., Böcker, S.: Coverage bias in small molecule machine learning. *Nature Communications* **16**, 554 (2025). <https://doi.org/10.1038/s41467-024-55462-w>
15. Lee, J., Lee, Y., Kim, J., Kosiorek, A., Choi, S., Teh, Y.W.: Set transformer: A framework for attention-based permutation-invariant neural networks. In: *Proceedings of the 36th International Conference on Machine Learning. Proceedings of Machine Learning Research*, vol. 97, pp. 3744–3753. PMLR (2019)
16. Liu, H., Bushuiev, R., Lightheart, I., Manjrekar, M., Bushuiev, A., Lederbauer, M., Jozefov, F., Wang, Y., Hassoun, S., Sivic, J., Taylor, J., Wang, R., Healey, D., Pluskal, T., Coley, C.W.: MassSpecGym in the wild: Uncovering and correcting evaluation pitfalls in AI-driven molecule discovery (2026). <https://doi.org/10.48550/arXiv.2606.19624>, preprint
17. Pang, L., Xu, J., Ai, Q., Lan, Y., Cheng, X., Wen, J.: SetRank: Learning a permutation-invariant ranking model for information retrieval. In: *Proceedings of the 43rd International ACM SIGIR Conference on Research and Development in Information Retrieval*. pp. 499–508 (2020). <https://doi.org/10.1145/3397271.3401104>
18. Rakhshaninejad, M., De Waele, G., Jürgens, M., Waegeman, W.: Reliable molecular retrieval from mass spectra using conformal prediction. *Journal of Chemical Information and Modeling* **66**(10), 5788–5800 (2026). <https://doi.org/10.1021/acs.jcim.6c00727>

19. Schymanski, E.L., Jeon, J., Gulde, R., Fenner, K., Ruff, M., Singer, H.P., Hollender, J.: Identifying small molecules via high resolution mass spectrometry: Communicating confidence. *Environmental Science & Technology* **48**(4), 2097–2098 (2014). <https://doi.org/10.1021/es5002105>
20. Stravs, M.A., Dührkop, K., Böcker, S., Zamboni, N.: MSNovelist: De novo structure generation from mass spectra. *Nature Methods* **19**(7), 865–870 (2022). <https://doi.org/10.1038/s41592-022-01486-3>
21. Tian, Y., Krishnan, D., Isola, P.: Contrastive multiview coding. In: *Computer Vision – ECCV 2020*. pp. 776–794. Springer (2020). https://doi.org/10.1007/978-3-030-58621-8_45
22. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L., Polosukhin, I.: Attention is all you need. In: *Advances in Neural Information Processing Systems*. vol. 30 (2017)
23. Xiong, P., Xu, H., Zheng, H.: Supervised contrastive learning leads to more reasonable spectral embeddings. *Analytical Chemistry* **97**(37), 20137–20146 (2025). <https://doi.org/10.1021/acs.analchem.5c02655>
24. Young, A., Röst, H., Wang, B.: Tandem mass spectrum prediction for small molecules using graph transformers. *Nature Machine Intelligence* **6**(4), 404–416 (2024). <https://doi.org/10.1038/s42256-024-00816-8>
25. Zhang, Y., Ding, K., Wu, Y., Zhuang, X., Yang, Y., Zhang, Q., Chen, H.: Breaking the modality barrier: Generative modeling for accurate molecule retrieval from mass spectra. In: *Proceedings of the Fortieth AAAI Conference on Artificial Intelligence*. vol. 40, pp. 1561–1569 (2026). <https://doi.org/10.1609/aaai.v40i2.37132>